


EEC 181 TPU Design Group

- Bradley Manzo
 - Computer Engineer
 - Data Orchestration, Benchmarking, Software Design, Presentation Design
- Austin York
 - Electrical Engineer
 - Hardware and Protocol Design
- Max Dhillon
 - Electrical Engineer
 - FIFO Design
- Kent Cassidy
 - Computer Engineer
 - Demo and Validation design



Goals of project

Assigned by Professor Ventakesh Akella

- 4096x4096 Matrix Multiplication
- Sparse MM Functionality
- Benchmark against software on comparable hardware
- Awarded Outstanding Senior Design for our class


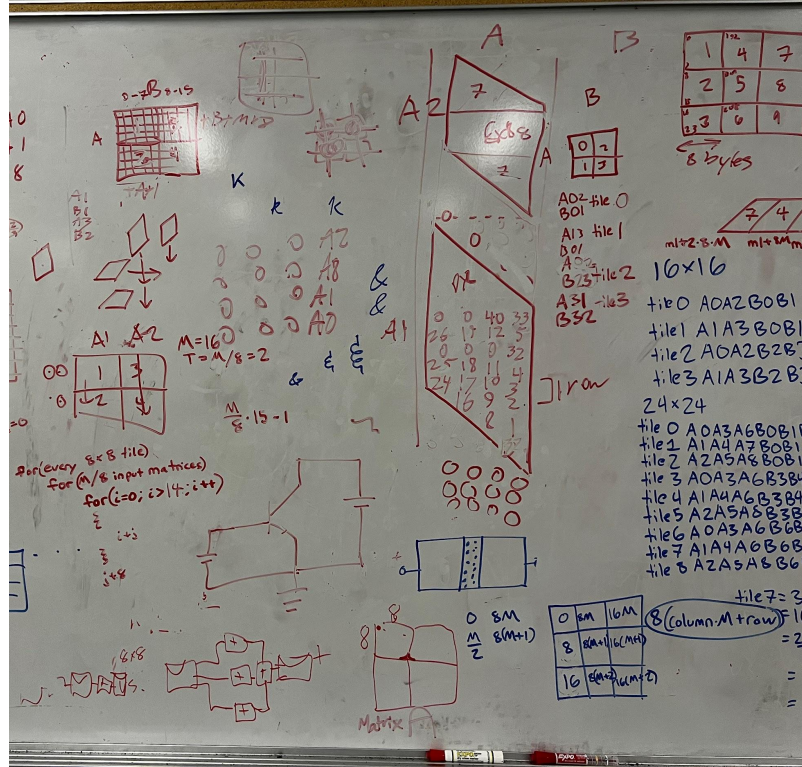
Resources

- Google TPU (2015)
- Cornell documentation on DMA for DE1-SoC

Challenges and Limitations

- Lack of experience/Knowledge in ML or DL Architecture/Design
- Limited to DE1-SoC Cyclone V FPGA
 - 32b bridge to communicate between HPS and FPGA
 - 128kB shared SRAM
 - Limited hardware, maximum Systolic array size: 8x8
 - Adder Carry Ripple Length

Teamwork (Software-Hardware Codesign)



Timeline

Winter Quarter


- Week 4 - Successfully wrote from program to shared SRAM
- Week 5 - FPGA operated on data sent to Shared SRAM
- Week 7 - Systolic Array working in Test bench
- Week 8 - Data Orchestration printing in correct order
- Week 9 - Single 8x8 functioning on TPU

Spring Quarter

- Week 6 - Completed Tiling Algorithm
- Week 7 - Implemented Tiling and added supporting commands for TPU
- Week 8 - Completed Sparse Algorithm
- Week 10 - Completed SRAM algorithm for $<4096 \times 4096$ MM


Results

Our Design vs Pure CPU over Matrix Size



Results


CPU Formatting/Writing/Reading Time vs Sparsity



*Percentage of zeros within input matrices

Results


TPU Time vs Sparsity



*Percentage of zeros within input matrices

Results

Memory Accesses Saved due to Sparsity



†Percentage of zeros within input matrices

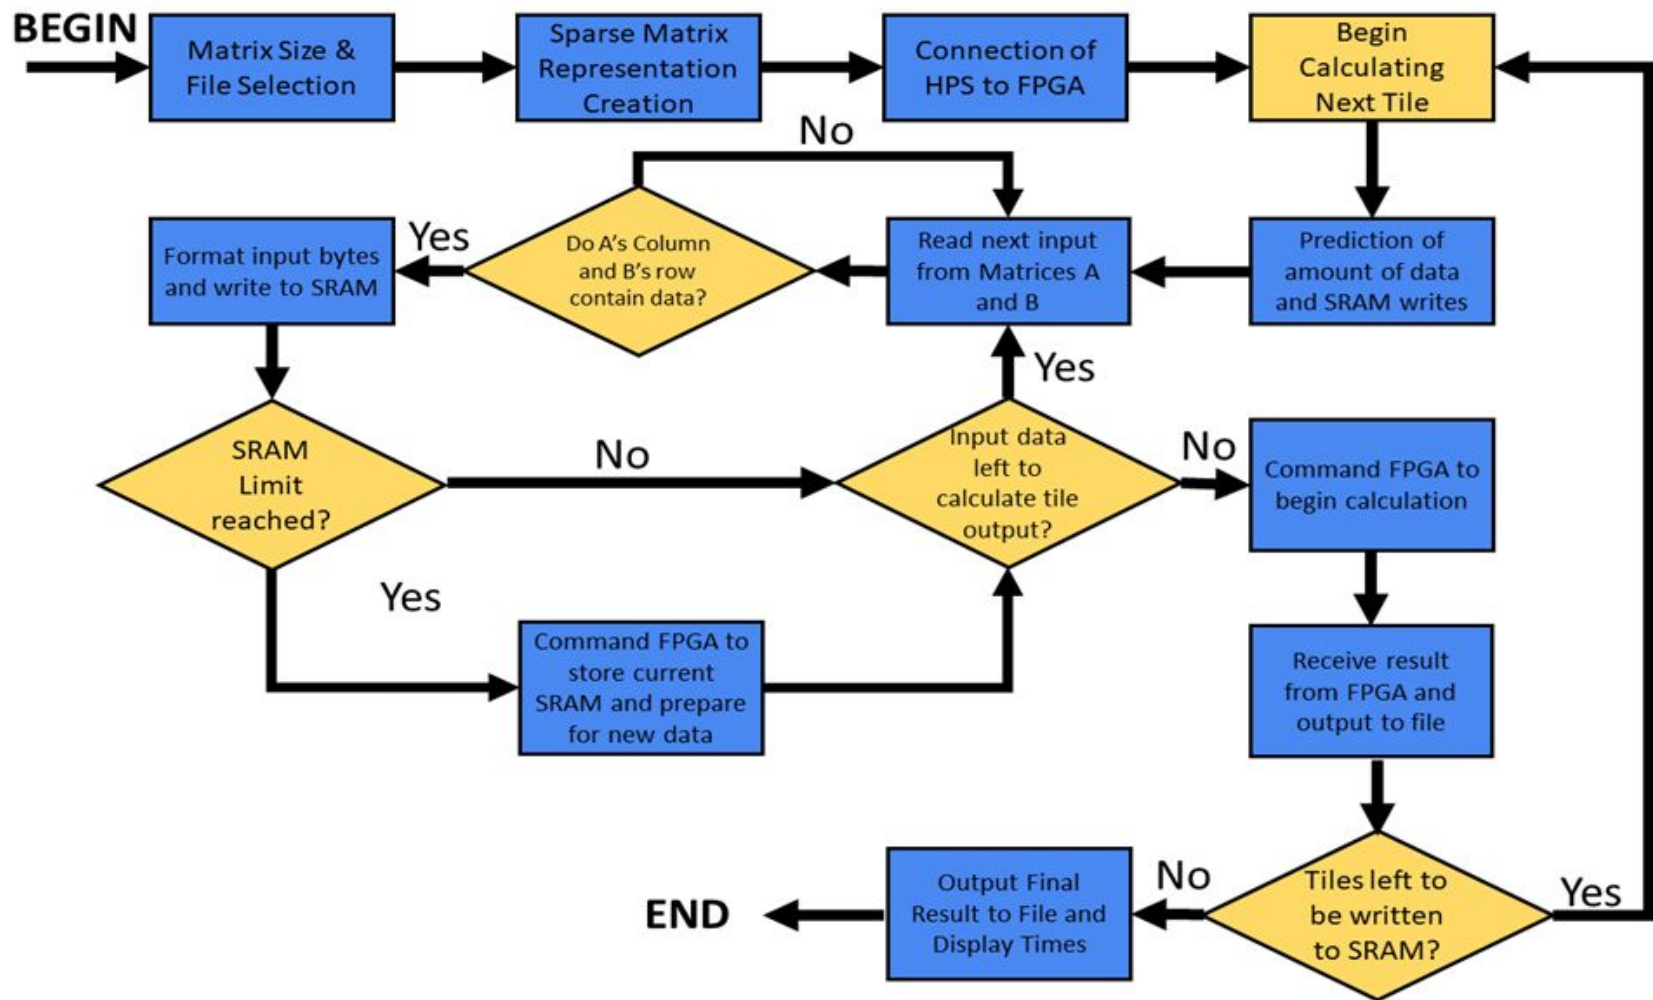
Future Goals

- Software Goals

- Convolution
- Parallel Threading
- Parameterize and Package Software as part of library function

- Hardware Goals

- FIFOs to pipeline HPS and TPU, instead of SRAM
- No data reuse -> Cache columns of Matrix B to reuse and reduce bandwidth
- **Store column of Matrix B in SRAM, FIFO Row of Matrix A**
- **Table Tile Partial Products and reuse instead of raw data**
- Accumulator Trees
- FPGA SDRAM Controller
- Decrease overall memory transitions using MEM blocks
- Reduce FPGA Logic and Interconnects



1. Accept Input Matrices from User

- User specifies dimensions (AH, AW, BH, BW)
 - If $AW \neq BH$, invalid inputs
 - $AH \times BW$ is the size of the output matrix
- Input width and height must be multiples of 8
 - 8x8 systolic array requires 64 elements per tile
 - Solution: pad partial tiles with zeros (maintains original structure)
 - `#define M_size(M) ((M % 8) ? (((M/8) + 1)*8) : (M))`
 - If a multiple of 8, no padding required
 - Otherwise, treat as next greatest multiple

1. Accept Input Matrices from User

AW = 8

0	8	16	24	32	40	48	56
1	9	17	25	33	41	49	57
2	10	18	26	34	42	50	58
3	11	19	27	35	43	51	59
4	12	20	28	36	44	52	60
5	13	21	29	37	45	53	61
6	14	22	30	38	46	54	62
7	15	23	31	39	47	55	63

AH = 8

M_size(AH) = 8

M_size(AW) = 8

1x1 Tile

AW = 9

0	9	18	27	36	45	54	63	72	0	0	0	0	0	0	0
1	10	19	28	37	46	55	64	73	0	0	0	0	0	0	0
2	11	20	29	38	47	56	65	74	0	0	0	0	0	0	0
3	12	21	30	39	48	57	66	75	0	0	0	0	0	0	0
4	13	22	31	40	49	58	67	76	0	0	0	0	0	0	0
5	14	23	32	41	50	59	68	77	0	0	0	0	0	0	0
6	15	24	33	42	51	60	69	78	0	0	0	0	0	0	0
7	16	25	34	43	52	61	70	79	0	0	0	0	0	0	0
8	17	26	35	44	53	62	71	80	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

AH = 9

M_size(AH) = 16

M_size(AW) = 16

2x2 Tiles

1. Accept Input Matrices from User

AW = 8

0	12	24	36	48	60	72	84
1	13	25	37	49	61	73	85
2	14	26	38	50	62	74	86
3	15	27	39	51	63	75	87
4	16	28	40	52	64	76	88
5	17	29	41	53	65	77	89
6	18	30	42	54	66	78	90
7	19	31	43	55	67	79	91
8	20	32	44	56	68	80	92
9	21	33	45	57	69	81	93
10	22	34	46	58	70	82	94
11	23	35	47	59	71	83	95
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

AH = 12

M_size(AH) = 16

M_size(AW) = 8

Tile Count = $(16/8) * (8/8) = 2$

AW = 10

0	8	16	24	32	40	48	56	64	72	0	0	0	0	0	0
1	9	17	25	33	41	49	57	65	73	0	0	0	0	0	0
2	10	18	26	34	42	50	58	66	74	0	0	0	0	0	0
3	11	19	27	35	43	51	59	67	75	0	0	0	0	0	0
4	12	20	28	36	44	52	60	68	76	0	0	0	0	0	0
5	13	21	29	37	45	53	61	69	77	0	0	0	0	0	0
6	14	22	30	38	46	54	62	70	78	0	0	0	0	0	0
7	15	23	31	39	47	55	63	71	79	0	0	0	0	0	0

AH = 8

M_size(AH) = 8

M_size(AW) = 16

1x2 Tiles

Tile Count = $(8/8) * (16/8) = 2$

1. Accept Input Matrices from User

- Process:
 - Input Matrices are:
 - File Stream (Volatile 1D array)
 - Column-Major Order
 - Binary Format (.bin)
 - Call M_size() to determine padded dimensions
 - Initialize a 1D array of type uint_8 of these dimensions with zeros
 - (M_size(W)*M_size(H)) long
 - Read each column of input matrix from file, to the base of corresponding column in Array

1. Accept Input Matrices from User

	0	1	2	3	4	5	6	7	8	9	10	11	...
...	12	13	14	15	16	17	18	19	20	21	22	23	...
...	24	25	26	27	28	29	30	31	32	33	34	35	...
...	36	37	38	39	40	41	42	43	44	45	46	47	...
...	48	49	50	51	52	53	54	55	56	57	58	59	...
...	60	61	62	63	64	65	66	67	68	69	70	71	...
...	72	73	74	75	76	77	78	79	80	81	82	83	...
...	84	85	86	87	88	89	90	91	92	93	94	95	...

Binary filestream

0	12	24	36	48	60	72	84
1	13	25	37	49	61	73	85
2	14	26	38	50	62	74	86
3	15	27	39	51	63	75	87
4	16	28	40	52	64	76	88
5	17	29	41	53	65	77	89
6	18	30	42	54	66	78	90
7	19	31	43	55	67	79	91
8	20	32	44	56	68	80	92
9	21	33	45	57	69	81	93
10	22	34	46	58	70	82	94
11	23	35	47	59	71	83	95
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Padded Tile
Visualization

	0	1	2	3	4	5	6	7	8	9	10	11	0	0	0	0	...
...	12	13	14	15	16	17	18	19	20	21	22	23	0	0	0	0	...
...	24	25	26	27	28	29	30	31	32	33	34	35	0	0	0	0	...
...	36	37	38	39	40	41	42	43	44	45	46	47	0	0	0	0	...
...	48	49	50	51	52	53	54	55	56	57	58	59	0	0	0	0	...
...	60	61	62	63	64	65	66	67	68	69	70	71	0	0	0	0	...
...	72	73	74	75	76	77	78	79	80	81	82	83	0	0	0	0	...
...	84	85	86	87	88	89	90	91	92	93	94	95	0	0	0	0	...


Local 1D uint_8 Array

2. Flag Sparse rows within input matrices

- If a row contains all zeros, result of that Systolic Array cycle will be zero
 - Can be ignored to reduce writes from SRAM
- Process:
 - Initialize boolean 1D arrays for Matrix A and B
 - 8 columns/rows per tile
 - Arrays of length $8 * \text{Tile Count}$
 - Arrays referenced before writing to SRAM

2. Flag Sparse rows within input matrices

- Matrix A scanned by column, since matrix A is written column by column to the Systolic Array
- A zero means the column isn't written to SRAM
- TPU behavior unaffected (Output tile is a result of exclusively non-zero inputs)



2. Flag Sparse rows within input matrices

- Matrix B scanned by row, since matrix B is written row by row to the Systolic Array
- A zero means the row isn't written to SRAM
- TPU behavior unaffected (Output tile is a result of exclusively non-zero inputs)

0	0	0	1	0	0	0	0
0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
1	0	0	1	1	0	0	0
0	0	0	0	0	0	0	0
0	1	1	0	0	0	1	0
0	0	0	0	0	1	0	0

1
1
0
0
1
0
1
1

3. Calculate Size of Input Data to write for each Output Tile

- Amount of input tiles required to calculate output
 - Determined by the common dimensions AW and BH

4x2 Tiles

A0	A4
A1	A5
A2	A6
A3	A7

2x3 Tiles

B0	B2	B4
B1	B3	B5

Output = 4x3 Tiles
Input Tiles per
Output Tile = 2

		B1		
		B1		
A4	A4	B1	C4	C8
A5	A5	C1	C5	C9
		C2	C6	C10
		C3	C7	C11

3. Calculate Size of Input Data to write for each Output Tile

- Ensure that Input Data to be written to SRAM wont overflow
 - 1024 Bytes (1KB) for each pair (A and B) of dense input tiles
 - SRAM region of the FPGA is of size 128 KB ($\frac{1}{8}$ MB!)
 - Limited to at most 128 pairs of Dense input tiles per Systolic Array Operation
 - Dense tile = 8 rows/columns => 1024 rows/columns

3. Calculate Size of Input Data to write for each Output Tile

- Sparsity introduces variable input sizes
 - Solution:
 - Simulate SRAM write, but with the Sparse Flag Matrices
 - If < 1024 rows/columns to be written, write all to SRAM, proceed normally
 - If > 1024 rows/columns
 - Write first 1024 rows/columns to TPU
 - Command TPU to calculate partial outputs
 - When TPU done with inputs, it will signal HPS to send next batch
 - When last batch written, TPU is commanded to write the output tile to SRAM

4. Write Input Data to SRAM

A03	A02	A01	A00	A Col 0
A07	A06	A05	A04	
B24	B16	B08	B00	B Row 0
B56	B48	B40	B32	
A11	A10	A09	A08	A Col 1
A15	A14	A13	A12	
B25	B17	B09	B01	B Row 1
B57	B49	B41	B33	
A19	A18	A17	A16	A Col 2
A23	A22	A21	A20	
B26	B18	B10	B02	B Row 2
B58	B50	B42	B34	
A27	A26	A25	A24	A Col 3
A31	A30	A29	A28	
B27	B19	B11	B03	B Row 3
B59	B51	B43	B35	
A35	A34	A33	A32	A Col 4
A39	A38	A37	A36	
B28	B20	B12	B04	B Row 4
B60	B52	B44	B36	
A43	A42	A41	A40	A Col 5
A47	A46	A45	A44	
B29	B21	B13	B05	B Row 5
B61	B53	B45	B37	
A51	A50	A49	A48	A Col 6
A55	A54	A53	A52	
B30	B22	B14	B06	B Row 6
B62	B54	B46	B38	
A59	A58	A57	A56	A Col 7
A63	A62	A61	A60	
B31	B23	B15	B07	B Row 7
B63	B55	B47	B39	

- Write to SRAM over 32b bridge
 - Limited to 4B or 4 inputs per write
 - Alternate complete col of A and row of B
- Process:
 - Program selects desired inputs from 1D arrays storing input matrices A and B
 - To read a column, +1 each read
 - To read a row, + Matrix Height each read

4. Write Input Data to SRAM

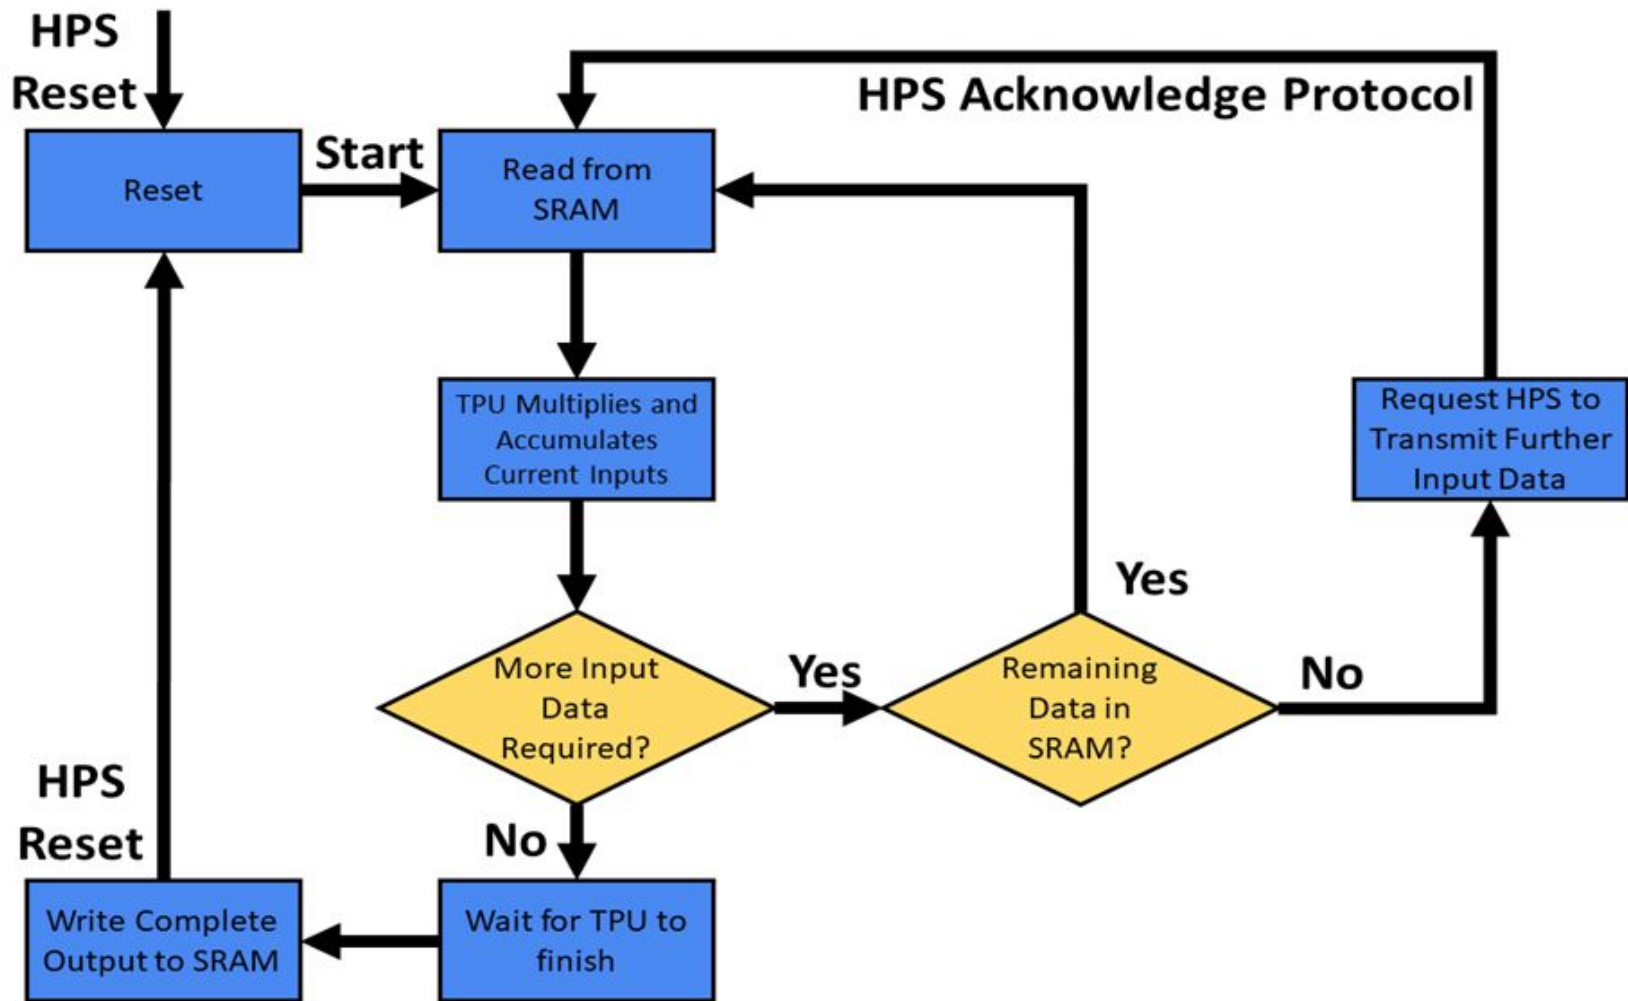
A03	A02	A01	A00	A Col 0
A07	A06	A05	A04	
B24	B16	B08	B00	B Row 0
B56	B48	B40	B32	
A11	A10	A09	A08	A Col 1
A15	A14	A13	A12	
B25	B17	B09	B01	B Row 1
B57	B49	B41	B33	
A19	A18	A17	A16	A Col 2
A23	A22	A21	A20	
B26	B18	B10	B02	B Row 2
B58	B50	B42	B34	
A27	A26	A25	A24	A Col 3
A31	A30	A29	A28	
B27	B19	B11	B03	B Row 3
B59	B51	B43	B35	
A35	A34	A33	A32	A Col 4
A39	A38	A37	A36	
B28	B20	B12	B04	B Row 4
B60	B52	B44	B36	
A43	A42	A41	A40	A Col 5
A47	A46	A45	A44	
B29	B21	B13	B05	B Row 5
B61	B53	B45	B37	
A51	A50	A49	A48	A Col 6
A55	A54	A53	A52	
B30	B22	B14	B06	B Row 6
B62	B54	B46	B38	
A59	A58	A57	A56	A Col 7
A63	A62	A61	A60	
B31	B23	B15	B07	B Row 7
B63	B55	B47	B39	

```
if(sparse_a[sparse_a_column] == 1 && sparse_b[sparse_b_row] == 1)
{
    //Write A Double Word
    for(byte = 0; byte < 4; byte++)
    {
        // Accumulates byte 03_02_01_00 (+3 reads down)
        SRAM_ACC_Write(&sram_ptr, &ACC, matrix_a[3 + a_offset - byte]);
    }
    for(byte = 0; byte < 4; byte++)
    {
        // Accumulates byte 07_06_05_04 (+7 reads down)
        SRAM_ACC_Write(&sram_ptr, &ACC, matrix_a[7 + a_offset - byte]);
    }
    // Write B Double Word
    for(byte = 0; byte < 4; byte++)
    {
        // Accumulates bytes 24_16_08_00 (increments column, by adding B height)
        SRAM_ACC_Write(&sram_ptr, &ACC, matrix_b[M_size(BH)*3 + b_offset + AB_DoubleWord - byte*M_size(BH)]);
    }
    for(byte = 0; byte < 4; byte++)
    {
        // Accumulates bytes 56_48_40_32 (increments column, by adding B height)
        SRAM_ACC_Write(&sram_ptr, &ACC, matrix_b[M_size(BH)*7 + b_offset + AB_DoubleWord - byte*M_size(BH)]);
    }
    sram_limit++;
    written_doublewords++;
}
```


4. Write Input Data to SRAM

A03	A02	A01	A00	A Col 0
A07	A06	A05	A04	
B24	B16	B08	B00	B Row 0
B56	B48	B40	B32	
A11	A10	A09	A08	A Col 1
A15	A14	A13	A12	
B25	B17	B09	B01	B Row 1
B57	B49	B41	B33	
A19	A18	A17	A16	A Col 2
A23	A22	A21	A20	
B26	B18	B10	B02	B Row 2
B58	B50	B42	B34	
A27	A26	A25	A24	A Col 3
A31	A30	A29	A28	
B27	B19	B11	B03	B Row 3
B59	B51	B43	B35	
A35	A34	A33	A32	A Col 4
A39	A38	A37	A36	
B28	B20	B12	B04	B Row 4
B60	B52	B44	B36	
A43	A42	A41	A40	A Col 5
A47	A46	A45	A44	
B29	B21	B13	B05	B Row 5
B61	B53	B45	B37	
A51	A50	A49	A48	A Col 6
A55	A54	A53	A52	
B30	B22	B14	B06	B Row 6
B62	B54	B46	B38	
A59	A58	A57	A56	A Col 7
A63	A62	A61	A60	
B31	B23	B15	B07	B Row 7
B63	B55	B47	B39	

- Write to SRAM over 32b bridge
 - Limited to 4B or 4 inputs per write
 - Alternate complete col of A and row of B
- Process:
 - Row and Column of input tile internally maintained
 - Beginning of tile from 0 = $8 * (column \cdot M_Height + row)$
 - `a_offset = 8 * (a_column * M_size(AH) + a_row);`
 - `b_offset = 8 * (b_column * M_size(BH) + b_row);`
 - Address of A column element = `Array[tile offset + i]`
 - Address of B row element = `Array[tile offset + i * M_Height]`
- Four 8b elements accumulated onto a 32b int, written to the next available address of SRAM



5. TPU Calculates Systolic Array Result



- Inputs written to SRAM in order
- TPU reads from SRAM
 - feeds each input into the top and left rows of MACS
- Outputs buffered when MAC completes, and written to SRAM when MAC 7,7 concludes
- **Output Stationary**

5. TPU Calculates Systolic Array Result

- MAC Specs
 - 1 8b Multiplier
 - 1 32b Adder
 - 64 23b registers for output stationary
- Performance
 - Amount of cycles to compute one tile
 - $(65 + 8 + (\text{TileWidth} * 8)) * (\text{TotalTiles})$
 - 65 cycles to write outputs
 - 8 cycles of overhead
 - 137 cycles per tile

6. Program receives Tile Output from TPU

- Outputs stored in column major order
- Written from SRAM to an output file
 - Local Array became too large for the ARM core
 - Zeros within partial tiles are neglected
- Tiles stored in column-major order
 - The 64 elements within each tile are also in column-major order
 - One tile to the tile below (+64)
 - One tile to the right (+64*Output_Matrix_Tile_Height)